

Q1

```
%pip install pandas
%pip install prophet
%pip install -U scikit-learn
```

```
import numpy as np
import pandas as pd
import prophet
import matplotlib.pyplot as plt
import scipy.optimize as opt
```

```
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (2.0.3)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas) (2023.4)
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-packages (from pandas) (2024.1)
Requirement already satisfied: numpy>=1.21.0 in /usr/local/lib/python3.10/dist-packages (from pandas) (1.25.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.2->pandas) (1.16.0)
Requirement already satisfied: prophet in /usr/local/lib/python3.10/dist-packages (1.1.5)
Requirement already satisfied: cmdstanpy>=1.0.4 in /usr/local/lib/python3.10/dist-packages (from prophet) (1.2.2)
Requirement already satisfied: numpy>=1.15.4 in /usr/local/lib/python3.10/dist-packages (from prophet) (1.25.2)
Requirement already satisfied: matplotlib>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from prophet) (3.7.1)
Requirement already satisfied: pandas>=1.0.4 in /usr/local/lib/python3.10/dist-packages (from prophet) (2.0.3)
Requirement already satisfied: holidays>=0.25 in /usr/local/lib/python3.10/dist-packages (from prophet) (0.49)
Requirement already satisfied: tqdm>=4.36.1 in /usr/local/lib/python3.10/dist-packages (from prophet) (4.66.4)
Requirement already satisfied: importlib-resources in /usr/local/lib/python3.10/dist-packages (from prophet) (6.4.0)
Requirement already satisfied: stanio<2.0.0,>=0.4.0 in /usr/local/lib/python3.10/dist-packages (from cmdstanpy>=1.0.4->prophet) (0.5.0)
Requirement already satisfied: python-dateutil in /usr/local/lib/python3.10/dist-packages (from holidays>=0.25->prophet) (2.8.2)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=2.0.0->prophet) (1.2.1)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=2.0.0->prophet) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=2.0.0->prophet) (4.51.0)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=2.0.0->prophet) (1.4.5)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=2.0.0->prophet) (24.0)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=2.0.0->prophet) (9.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=2.0.0->prophet) (3.1.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=1.0.4->prophet) (2023.4)
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=1.0.4->prophet) (2024.1)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil->holidays>=0.25->prophet) (1.16.0)
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.10/dist-packages (1.2.2)
Collecting scikit-learn
  Downloading scikit_learn-1.5.0-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (13.3 MB)
13.3/13.3 MB 58.5 MB/s eta 0:00:00
Requirement already satisfied: numpy>=1.19.5 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.25.2)
Requirement already satisfied: scipy>=1.6.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.11.4)
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.4.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (3.5.0)
Installing collected packages: scikit-learn
  Attempting uninstall: scikit-learn
    Found existing installation: scikit-learn 1.2.2
    Uninstalling scikit-learn-1.2.2:
      Successfully uninstalled scikit-learn-1.2.2
  Successfully installed scikit-learn-1.5.0
```

#Regression function

```
from sklearn import linear_model
reg = linear_model.LinearRegression()

#data = pd.read_csv('/SmallBusiness.csv')
#years = data['Year']
#sales = data['Sales']

years = np.array((range(1,21)))
sales = np.array ([283, 288, 336, 388, 406, 412, 416, 435, 428, 435, 462, 452, 474, 476, 497, 487, 523, 528, 532, 552])


df = pd.DataFrame.from_dict({'Year':years,'Sales':sales})
print(df)

reg = reg.fit(df[['Year']], df[['Sales']])
print(' intercept:', reg.intercept_)#b0
print(' coefficient:', reg.coef_)#b1
print(reg.score(df[['Year']], df[['Sales']]))#R^2
```

	Year	Sales
0	1	283
1	2	288
2	3	336
3	4	388
4	5	406
5	6	412
6	7	416
7	8	435
8	9	428

```
9      10      435
10     11      462
11     12      452
12     13      474
13     14      476
14     15      497
15     16      487
16     17      523
17     18      528
18     19      532
19     20      552
intercept: [312.1]
coefficient: [[12.22857143]]
0.9194295594103281



#calculate the prediction (use the function above)
df['Prediction']=reg.predict(df[['Year']])
#calculate error (actual value - predict value)
df['Error']=df['Sales']-df['Prediction']
print(df)
```



	Year	Sales	Prediction	Error
0	1	283	324.328571	-41.328571
1	2	288	336.557143	-48.557143
2	3	336	348.785714	-12.785714
3	4	388	361.014286	26.985714
4	5	406	373.242857	32.757143
5	6	412	385.471429	26.528571
6	7	416	397.700000	18.300000
7	8	435	409.928571	25.071429
8	9	428	422.157143	5.842857
9	10	435	434.385714	0.614286
10	11	462	446.614286	15.385714
11	12	452	458.842857	-6.842857
12	13	474	471.071429	2.928571
13	14	476	483.300000	-7.300000
14	15	497	495.528571	1.471429
15	16	487	507.757143	-20.757143
16	17	523	519.985714	3.014286
17	18	528	532.214286	-4.214286
18	19	532	544.442857	-12.442857
19	20	552	556.671429	-4.671429

```
plt.plot (df['Year'], df['Sales'], "o", label="data")
plt.plot (df['Year'], df['Prediction'], "b-", label="OLS") #Ordinary Least Squares
plt.legend (loc="best")
plt.xlabel ('Year')
plt.ylabel ('Sales')
plt.show ()

plt.plot (df['Prediction'], df['Error'], "o")
plt.xlabel ("Prediction")
plt.ylabel ("Residual") #The errors
plt.show ()
```



```
#Exponential smoothing
#from scipy.optimize import minimize

def exp(sales, extra_periods=1, alpha = 0.5):

    cols = len(sales)
    sales = np.append(sales, [np.nan]*extra_periods)

    f = np.full(cols+extra_periods, np.nan)
    f[1] = sales[0]

    for t in range(2, cols+1):
        f[t] = alpha*sales[t-1] + (1-alpha)*f[t-1]

    # forecast
    for t in range(cols+1, cols+extra_periods):
        f[t] = f[t-1]

    df = pd.DataFrame.from_dict({'Sales':sales, 'Forecast':f, 'ES_Error':sales-f})

    return df

def exp_smooth_opti(sales, extra_periods = 6):

    params = [] # contains all the different parameter sets
    KPIs = [] # contains the results of each model
    dfs = [] # contains all the dataframes returned by the different models

    for alpha in [0.05, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9]:

        df = exp(sales, extra_periods=extra_periods, alpha=alpha)
        params.append(f'Simple Smoothing, alpha: {alpha}')
        dfs.append(df)
        MAE = df['ES_Error'].abs().mean()
        KPIs.append(MAE)
```

```
mini = np.argmin(KPIs)
print(f'Best solution found for {params[mini]} MAE of', round(KPIs[mini], 2))

return dfs[mini]

df2 = exp_smooth_opti(sales)

print (df2)
```

↻ Best solution found for Simple Smoothing, alpha: 0.9 MAE of 17.73

	Sales	Forecast	ES_Error
0	283.0	NaN	NaN
1	288.0	283.000000	5.000000
2	336.0	287.500000	48.500000
3	388.0	331.150000	56.850000
4	406.0	382.315000	23.685000
5	412.0	403.631500	8.368500
6	416.0	411.163150	4.836850
7	435.0	415.516315	19.483685
8	428.0	433.051631	-5.051631
9	435.0	428.505163	6.494837
10	462.0	434.350516	27.649484
11	452.0	459.235052	-7.235052
12	474.0	452.723505	21.276495
13	476.0	471.872351	4.127649
14	497.0	475.587235	21.412765
15	487.0	494.858724	-7.858724
16	523.0	487.785872	35.214128
17	528.0	519.478587	8.521413
18	532.0	527.147859	4.852141
19	552.0	531.514786	20.485214
20	NaN	549.951479	NaN
21	NaN	549.951479	NaN
22	NaN	549.951479	NaN
23	NaN	549.951479	NaN
24	NaN	549.951479	NaN
25	NaN	549.951479	NaN

```
SSE_reg = (df['Error']**2).sum()
SSE_ES = (df2['ES_Error']**2).sum()

print(SSE_reg)
print(SSE_ES)
```

↻

8714.257142857146
10273.563718144005

Answer

(a) According to the above code and result, b0 (intercept = 312.1) and b1 (coefficient = 12.22857143), so the estimated regression function will be : $\hat{Y} = 312.1 + 12.22857143 \cdot x$

(b) R^2 is the value which we use to measure how well the regression line approximate the real data in statistical. It ranges from 0 to 1. If the value are close to 1, which indicates that this regression line can fits the data better. Otherwise, the regression can not explain the variables of the response data around its mean at all. In this case, the value of R^2 is 0.9194295594103281 (according to the above code). It means that about 91.94% of the variance in this data can be explained by the time value (x:years).

(c) First, identify the value of alpha (smoothing constant). We use the optimal alpha (0.9), which we found in the code. Then, adopt the first sales value in actual data as the initial forecast (F1). For each subsequent period, the function is : $F_t = \alpha A_{t-1} + (1-\alpha)F_{t-1}$, use this function to calculate prediction in each year (like above code and results). For example, the next prediction of F1 is 283, and the second one is 287.5. In linear trend model, we can captures a consistent trend over time. It's useful when data show a linear pattern, and also can see the relation within each points. If we use expoential smoothing model, we can adapt quickly to new trend. In this case, SSE of linear trend model is smaller then expoential smoothing model, which means the prediction is closer to its actual value. As a result, we can consider to adopt linear trend model.

2-1

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

data = {
    'Quarter': [
        'Q1-1985', 'Q2-1985', 'Q3-1985', 'Q4-1985',
        'Q1-1986', 'Q2-1986', 'Q3-1986', 'Q4-1986',
        'Q1-1987', 'Q2-1987', 'Q3-1987', 'Q4-1987',
        'Q1-1988', 'Q2-1988', 'Q3-1988', 'Q4-1988',
        'Q1-1989', 'Q2-1989', 'Q3-1989', 'Q4-1989'
    ],
    'Shipments': [
        4009, 4321, 4224, 3944,
        4123, 4522, 4657, 4030,
        4493, 4806, 4551, 4485,
        4595, 4799, 4417, 4258,
        4245, 4900, 4585, 4533
    ]
}

df = pd.DataFrame(data)


def moving_average(d, extra_periods=1, n=3):
    cols = len(d)
    d = np.append(d, [np.nan]*extra_periods)
    f = np.full(cols+extra_periods, np.nan)
    for t in range(n, cols):
        f[t] = np.mean(d[t-n:t])
    f[t+1:] = np.mean(d[t-n+1:t+1])
    df = pd.DataFrame.from_dict({'Demand':d, 'Forecast':f, 'Error':d-f})
    return df

result = moving_average(df['Shipments'].values, extra_periods=1, n=4)

print(result)

result[['Demand', 'Forecast']].plot()
plt.title('Demand vs Forecast')
plt.xlabel('Time')
plt.ylabel('Shipments')
plt.grid(True)
plt.show()
```

	Demand	Forecast	Error
0	4009.0	NaN	NaN
1	4321.0	NaN	NaN
2	4224.0	NaN	NaN
3	3944.0	NaN	NaN
4	4123.0	4124.50	-1.50
5	4522.0	4153.00	369.00
6	4657.0	4203.25	453.75
7	4030.0	4311.50	-281.50
8	4493.0	4333.00	160.00
9	4806.0	4425.50	380.50
10	4551.0	4496.50	54.50
11	4485.0	4470.00	15.00
12	4595.0	4583.75	11.25
13	4799.0	4609.25	189.75
14	4417.0	4607.50	-190.50
15	4258.0	4574.00	-316.00
16	4245.0	4517.25	-272.25
17	4900.0	4429.75	470.25
18	4585.0	4455.00	130.00
19	4533.0	4497.00	36.00
20	NaN	4565.75	NaN



在1989 Q4 shipment有下降的趨勢，但我們計算出的shipment卻有上升的徵兆，因此有可能會overestimate 1990 Q1的shipment，理由主要是因為利用moving average的算法算出的預測值通常無法及時反映當下的數值，像是1989 Q4 shipment有下降的趨勢，但我們的預測可能要再晚一期的時間才能反應過來。

2-2

```
import pandas as pd
import statsmodels.api as sm

data = {
    'Year': ['1985', '1985', '1985', '1985', '1986', '1986', '1986', '1986', '1987', '1987', '1987', '1987', '1988', '1988', '1988', '1988',
    'Quarter': ['Q1', 'Q2', 'Q3', 'Q4', 'Q1', 'Q2', 'Q3', 'Q4', 'Q1', 'Q2', 'Q3', 'Q4', 'Q1', 'Q2', 'Q3', 'Q4', 'Q1', 'Q2', 'Q3', 'Q4'],
    'Shipments': [4009, 4321, 4224, 3944, 4123, 4522, 4657, 4030, 4493, 4806, 4551, 4485, 4595, 4799, 4417, 4258, 4245, 4900, 4585, 4533]
}
df = pd.DataFrame(data)
df['Year'] = df['Year'].astype(int)
df['Time'] = range(1, len(df) + 1)

seasonal_dummies = pd.get_dummies(df['Quarter']).astype(int)
df = pd.concat([df, seasonal_dummies], axis=1)

train = df[df['Year'] != 1989]
test = df[df['Year'] == 1989]
X_train = sm.add_constant(train[['Time', 'Q1', 'Q2', 'Q3']])
y_train = train['Shipments']
X_test = sm.add_constant(test[['Time', 'Q1', 'Q2', 'Q3']])
y_test = test['Shipments']

model = sm.OLS(y_train, X_train).fit()

predictions = model.predict(X_test)
print(predictions)

print(model.summary())
```

16 4662.25
17 4969.25
18 4819.50
19 4536.50
dtype: float64

OLS Regression Results

=====

Dep. Variable:	Shipments	R-squared:	0.749
Model:	OLS	Adj. R-squared:	0.658
Method:	Least Squares	F-statistic:	8.208
Date:	Tue, 28 May 2024	Prob (F-statistic):	0.00255
Time:	15:04:56	Log-Likelihood:	-100.84
No. Observations:	16	AIC:	211.7
Df Residuals:	11	BIC:	215.5
Df Model:	4		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	3822.0000	119.462	31.994	0.000	3559.067	4084.933
Time	35.7250	8.904	4.012	0.002	16.127	55.323
Q1	232.9250	115.754	2.012	0.069	-21.848	487.698
Q2	504.2000	114.029	4.422	0.001	253.224	755.176
Q3	318.7250	112.981	2.821	0.017	70.055	567.395

Omnibus:	0.690	Durbin-Watson:	1.504
Prob(Omnibus):	0.708	Jarque-Bera (JB):	0.345
Skew:	0.345	Prob(JB):	0.841
Kurtosis:	2.799	Cond. No.	45.9

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

/usr/local/lib/python3.10/dist-packages/scipy/stats/_stats_py.py:1806: UserWarning: kurtosistest only valid for n>=20 ... continuing anyway: warnings.warn("kurtosistest only valid for n>=20 ... continuing ")

df



	Year	Quarter	Shipments	Time	Q1	Q2	Q3	Q4
0	1985	Q1	4009	1	1	0	0	0
1	1985	Q2	4321	2	0	1	0	0
2	1985	Q3	4224	3	0	0	1	0
3	1985	Q4	3944	4	0	0	0	1
4	1986	Q1	4123	5	1	0	0	0
5	1986	Q2	4522	6	0	1	0	0
6	1986	Q3	4657	7	0	0	1	0
7	1986	Q4	4030	8	0	0	0	1
8	1987	Q1	4493	9	1	0	0	0
9	1987	Q2	4806	10	0	1	0	0
10	1987	Q3	4551	11	0	0	1	0
11	1987	Q4	4485	12	0	0	0	1
12	1988	Q1	4595	13	1	0	0	0
13	1988	Q2	4799	14	0	1	0	0
14	1988	Q3	4417	15	0	0	1	0
15	1988	Q4	4258	16	0	0	0	1
16	1989	Q1	4245	17	1	0	0	0
17	1989	Q2	4900	18	0	1	0	0
18	1989	Q3	4585	19	0	0	1	0
19	1989	Q4	4533	20	0	0	0	1

```
import numpy as np
import pandas as pd

def holt_winters_multiplicative(d, slen=4, extra_periods=1, alpha=0.4, beta=0.4, gamma=0.3):
    cols = len(d)
    d = np.append(d, [np.nan] * extra_periods)

    f, E, T, S = np.full((4, cols + extra_periods), np.nan)

    for i in range(slen):
        E[i] = d[:slen].mean() # Initial average
        T[i] = 0 # Initial trend
        S[i] = d[i] / E[i] # Initial seasonality index

    # t+1 forecast
    for t in range(slen, cols):
        f[t] = (E[t-1] + T[t-1]) * S[t-slen]
        E[t] = alpha * (d[t] / S[t-slen]) + (1-alpha) * (E[t-1] + T[t-1])
        T[t] = beta * (E[t] - E[t-1]) + (1-beta) * T[t-1]
        S[t] = gamma * (d[t] / E[t]) + (1-gamma) * S[t-slen]

    for t in range(cols, cols + extra_periods):
        f[t] = (E[cols-1] + (t-(cols-1)) * T[cols-1]) * S[t-slen]

    df = pd.DataFrame.from_dict({'Demand': d, 'Forecast': f, 'Level': E, 'Trend': T, 'Season': S, 'Error': d - f})

    return df
```

```
shipments = data['Shipments']
print(type(shipments))
result = holt_winters_multiplicative(shipments, slen=4, extra_periods=4, alpha=0.4, beta=0.4, gamma=0.3)
print(result[['Demand', 'Forecast', 'Level', 'Trend', 'Season', 'Error']])
```

	<class 'list'>					
	Demand	Forecast	Level	Trend	Season	Error
0	4009.0	NaN	4124.500000	0.000000	0.971997	NaN
1	4321.0	NaN	4124.500000	0.000000	1.047642	NaN
2	4224.0	NaN	4124.500000	0.000000	1.024124	NaN
3	3944.0	NaN	4124.500000	0.000000	0.956237	NaN
4	4123.0	4009.000000	4171.413744	18.765498	0.976916	114.000000
5	4522.0	4389.808341	4240.651308	38.954324	1.053253	132.191659
6	4657.0	4382.847422	4386.683502	81.785472	1.035374	274.152578
7	4030.0	4272.915902	4366.855726	41.140173	0.946224	-242.915902
8	4493.0	4306.240845	4484.464783	71.727727	0.984412	186.759155
9	4806.0	4798.824280	4558.917674	72.817792	1.053537	7.175720
10	4551.0	4795.576231	4537.247351	35.022546	1.025671	-244.576231
11	4485.0	4326.392544	4639.318465	61.841975	0.952378	158.607456
12	4595.0	4627.878973	4687.800601	56.498038	0.983150	-32.878973
13	4799.0	4998.291812	4668.632799	26.231702	1.045853	-199.291812
14	4417.0	4815.385322	4539.498773	-35.914606	1.009874	-398.385322
15	4258.0	4289.114467	4490.516010	-41.141852	0.951131	-31.114467
16	4245.0	4374.400312	4396.726905	-62.200754	0.977852	-129.400312
17	4900.0	4533.276196	4474.784441	-6.097438	1.060604	366.723804
18	4585.0	4512.810909	4497.280309	5.339885	1.012763	72.189091
19	4533.0	4282.580737	4607.934526	47.465617	0.960913	250.419263
20	NaN	4552.292271	NaN	NaN	NaN	NaN
21	NaN	4987.880050	NaN	NaN	NaN	NaN
22	NaN	4810.961426	NaN	NaN	NaN	NaN
23	NaN	4610.265214	NaN	NaN	NaN	NaN

✓ 2-3

✓ 2-3

```
[ ] result_for_hw_model = result['Forecast'][-4:].values
validation = shipments[16:]
hw_mape = np.mean(np.abs((np.array(validation) - result_for_hw_model) / np.array(validation))) * 100
print(f'Holt-Winters MAPE: {hw_mape:.2f}%')
```

⇒ Holt-Winters MAPE: 3.92%

```
[ ] linear_mape = np.mean(np.abs((y_test.values - predictions) / y_test.values)) * 100
print(f'Linear Trend and Seasonal Model MAPE: {linear_mape:.2f}%')
```

⇒ Linear Trend and Seasonal Model MAPE: 4.11%

```
▶ if hw_mape < linear_mape:
    print("Holt-Winters 模型更適合預測 Q1-1990 的 shipment")
else:
    print("線性模型更適合預測 Q1-1990 的 shipment")
```

⇒ Holt-Winters 模型更適合預測 Q1-1990 的 shipment

```
import numpy as np
import pandas as pd
import prophet
from google.colab import files
import io
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
import matplotlib.pyplot as plt
import statsmodels.api as sm
import scipy.stats as stats
```

```
# Upload and read the CSV file
uploaded = files.upload()
```



選擇檔案 Duque.csv

- **Duque.csv**(text/csv) - 400 bytes, last modified: 2024/5/29 - 100% done
- Saving Duque.csv to Duque.csv

```
df1 = pd.read_csv(io.BytesIO(uploaded['Duque.csv']))
```

```
# Extract temperature and demand data
demand_data = df1.iloc[:, 3].values
temperature_data=df1.iloc[:, 2].values.reshape(-1, 1)
```

✓ Q3(a)

```
X_train, X_test, y_train, y_test = train_test_split(temperature_data, demand_data, test_size=0.2, random_state=0)
simple_regressor = LinearRegression()
simple_regressor.fit(X_train, y_train)
y_pred_simple = simple_regressor.predict(temperature_data)
```

```
slope = simple_regressor.coef_[0]
```


```
intercept = simple_regressor.intercept_
```

```
print(f'Estimated regression equation: y = {slope:.4f} * x + {intercept:.4f}')
```

```
plt.scatter(temperature_data, demand_data, color='black', label='Actual data')
plt.plot(temperature_data, y_pred_simple, color='blue', linewidth=3, label='Regression line')
plt.xlabel('temperature')
plt.ylabel('demand')
plt.legend()
plt.show()
```



Estimated regression equation: $y = 0.3962 * x + 28.0319$



✓ Q3(b)

```
# Create dummy variables for the day of the week
df2 = pd.get_dummies(df1, columns=['Day'], drop_first=True)

# Define feature columns for multiple regression
feature_columns = ['Temp'] + [col for col in df2.columns if col.startswith('Day_')]

# Ensure feature columns are correct
print("Feature columns: ", feature_columns)

# Prepare the data for multiple linear regression
X = df2[feature_columns].values
y = df2['Demand'].values

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)

# Create and train the multiple regression model
multiple_regressor = LinearRegression()
multiple_regressor.fit(X_train, y_train)

# Estimated regression model
coefficients = multiple_regressor.coef_
intercept = multiple_regressor.intercept_
equation = f"y = {intercept:.4f} + " + " + " + ".join([f"{coeff:.4f}*{name}" for coeff, name in zip(coefficients, feature_columns)])
print('Estimated regression equation:', equation)

# Interpret the coefficients
for feature, coeff in zip(feature_columns, coefficients):
    print(f"Coefficient for {feature}: {coeff:.4f}")
```

Feature columns: ['Temp', 'Day_2', 'Day_3', 'Day_4', 'Day_5', 'Day_6', 'Day_7']
 Estimated regression equation: $y = 5.0303 + 0.5859 \cdot \text{Temp} + 7.2803 \cdot \text{Day}_2 + 9.0474 \cdot \text{Day}_3 + 8.4094 \cdot \text{Day}_4 + 12.0423 \cdot \text{Day}_5 + 4.8906 \cdot \text{Day}_6 + -2.6914 \cdot \text{Day}_7$
 Coefficient for Temp: 0.5859
 Coefficient for Day_2: 7.2803
 Coefficient for Day_3: 9.0474
 Coefficient for Day_4: 8.4094
 Coefficient for Day_5: 12.0423
 Coefficient for Day_6: 4.8906
 Coefficient for Day_7: -2.6914

✓ Q3(c)

```
y_pred_multiple = multiple_regressor.predict(X)

# Plot the actual vs predicted values for simple regression
plt.figure(figsize=(10, 6))
plt.plot(df1.index, demand_data, color='black', label='Actual data')
plt.plot(df1.index, y_pred_simple, color='blue', label='Predicted data (Simple Regression)', alpha=0.7)
plt.xlabel('Day')
plt.ylabel('Demand')
plt.legend()
plt.title('Actual vs Predicted Peak Demand (Simple Regression)')
plt.show()

# Plot the actual vs predicted values for multiple regression
plt.figure(figsize=(10, 6))
plt.plot(df1.index, demand_data, color='black', label='Actual data')
plt.plot(df1.index, y_pred_multiple, color='red', label='Predicted data (Multiple Regression)', alpha=0.7)
plt.xlabel('Day')
plt.ylabel('Demand')
plt.legend()
plt.title('Actual vs Predicted Peak Demand (Multiple Regression)')
plt.show()

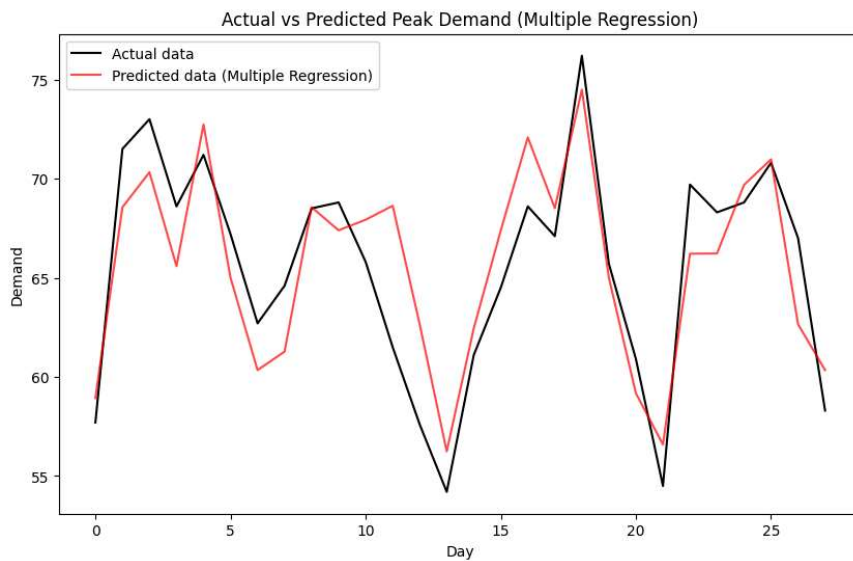

# Evaluate model performance
from sklearn.metrics import mean_absolute_error, mean_squared_error

mae_simple = mean_absolute_error(demand_data, y_pred_simple)
mae_multiple = mean_absolute_error(demand_data, y_pred_multiple)

mse_simple = mean_squared_error(demand_data, y_pred_simple)
mse_multiple = mean_squared_error(demand_data, y_pred_multiple)

rmse_simple = np.sqrt(mse_simple)
rmse_multiple = np.sqrt(mse_multiple)

print(f'Simple Model - MAE: {mae_simple}, MSE: {mse_simple}, RMSE: {rmse_simple}')
print(f'Multiple Model - MAE: {mae_multiple}, MSE: {mse_multiple}, RMSE: {rmse_multiple}')
```



Simple Model - MAE: 4.598403138048843, MSE: 28.61984969829599, RMSE: 5.34975230251793
 Multiple Model - MAE: 2.338567275094857, MSE: 7.596797956025045, RMSE: 2.756228937526

✓ Q3(d)

```
#跑出回歸式
df_with_dummies = pd.get_dummies(df1, columns=['Day'], drop_first=True)

feature_columns = ['Temp'] + [col for col in df_with_dummies.columns if col.startswith('Day_')]

print("Feature columns: ", feature_columns)

X = df_with_dummies[feature_columns].values
y = df_with_dummies['Demand'].values

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)

multiple_regressor = LinearRegression()
multiple_regressor.fit(X_train, y_train)
```

```
coefficients = multiple_regressor.coef_  
intercept = multiple_regressor.intercept_  
equation = f"y = {intercept:.4f} + " + " + ".join([f"{coeff:.4f}*{name}" for coeff, name in zip(coefficients, feature_columns)])  
print('Estimated regression equation:', equation)  
  
#進行新預測  
X_new = np.array([[94, 0, 1, 0, 0, 0, 0]]) #溫度94, Day_3的預測  
y_pred = multiple_regressor.predict(X_new)[0]  
residuals = y_train - multiple_regressor.predict(X_train)  
sigma = np.sqrt(np.sum(residuals**2) / (len(y_train) - len(coefficients) - 1))  
  
X_new_with_intercept = np.insert(X_new, 0, 1, axis=1)  
X_train_with_intercept = np.insert(X_train, 0, 1, axis=1)  
X_train_with_intercept = X_train_with_intercept.astype('float64')  
se_pred = np.sqrt(sigma**2 + X_new_with_intercept @ np.linalg.inv(X_train_with_intercept.T @ X_train_with_intercept) @ X_new_with_intercept.T  
  
t_value = stats.t.ppf(0.975, df=len(y_train) - len(coefficients) - 1)  
prediction_interval = (y_pred - t_value * se_pred, y_pred + t_value * se_pred)  
  
print(f"預測值: {y_pred:.2f}")  
  
print("範圍: ", prediction_interval[0][0], "到", prediction_interval[1][0])  
  
↗ Feature columns: ['Temp', 'Day_2', 'Day_3', 'Day_4', 'Day_5', 'Day_6', 'Day_7']  
Estimated regression equation: y = 5.0303 + 0.5859*Temp + 7.2803*Day_2 + 9.0474*Day_3 + 8.4094*Day_4 + 12.0423*Day_5 + 4.8906*Day_6 + -2.6914*Day_7  
預測值: 69.15  
範圍: [62.4041984] 到 [75.90094069]
```

有95%的機率，用電高峰會落在62.404到75.9之間(正負一個標準差)。